

Find Target Data Fast! A Method and Its Behavior of Target Data Collection using Online Machine Learning

Ruide LI[†], Yoko YAMAKATA^{††}, and Keishi TAJIMA[†]

[†] Kyoto University

36-1 Yoshida-Honmachi, Sakyo-ku, Kyoto 606-8501 Japan

^{††} The University of Tokyo

7-3-1 Hongo, Bunkyo-ku, Tokyo 113-8656 Japan

E-mail: †li.ruide@dl.soc.i.kyoto-u.ac.jp, ††yamakata@hal.t.u-tokyo.ac.jp, †††tajima@i.kyoto-u.ac.jp

Abstract In this paper, we discuss a problem of collecting data of a target class from a fixed pool of unlabeled dataset as fast as possible. Our method consists of two phases. In the first phase, we repeatedly choose a candidate data and query human annotators for the label. During that, we use the obtained labeled data also for training a classifier, and use it to choose the next candidate. When the classifier has been trained enough, we switch to the second phase. We apply the obtained classifier to the remaining data to collect target data. In the first phase, we want to choose candidates which are most likely to be target data in order to collect them fast. For the second phase, however, we want to query a label for a data which is likely to improve the classifier. We have a dilemma between these two. Another issue is difficulty of the estimation of the accuracy of the obtained classifier. The labeled data set obtained in the first phase is biased towards the target class, and by using that we need to estimate the accuracy of the current classifier for the remaining data which is biased in the opposite way. In this paper, we explain the difficulty of this new problem by using our experimental results, and also show the behavior of our simple method of ballancing the dilemma explained above.

Key words active learning, annotation, crowdsourcing, learn-to-enumerate

1 Introduction

In the pool-based active learning problem [1], we usually have a pool of unlabeled data, and the learner pick up a sample from the pool for human annotation, referring to a function which gives the priority of every sample. In this paper, we discuss a slightly different problem: a problem of collecting data of a target class from a fixed pool of unlabeled dataset as fast as possible by using machine learning but without any prelabeled training data. We propose a method which consists of two phases. In the first phase, we choose a candidate data one by one, and repeatedly query human annotators for the label. If it is labeled as the target class, we store it as a target data. During the first phase, we use the obtained labeled data also for train a classifier, and use the classifier to choose the next candidate. When we determine that the classifier has been trained enough, we switch to the second phase: we switch from human annotation to prediction by the classifier. We apply the obtained classifier to the remaining unlabelled data and collect data predicted as the target class.

This problem is different from the ordinary active learning

problem in several ways. First, our goal is to collect target data as fast as possible (i.e., with as few queries for non-target data). In the first phase, therefore, we want to choose a candidate which is most likely to be of the target class from the remaining unlabelled data. The performance of the obtained classifier is not important as long as we can collect target data fast. Moreover, the classification boundary of the classifier is not important in our problem as long as the single best candidate chosen by the classifier in each step is always a target data. On the other hand, in the ordinary active learning, we want to achieve high accuracy of the classifier while minimizing the total number of both target and non-target queries. In the ordinary active learning, therefore, uncertainty sampling [2], in which samples with lower confidence will be annotated earlier, is a well-known approach. Existing research has shown that the opposite approach, which gives higher priority to samples with higher confidence, is a better strategy [3] if we simply want to collect target data as fast as possible.

For the second phase, however, we also want to achieve high accuracy of the classifier. As a result, we have a dilemma between querying a label for a data which is likely

to be a target data and querying for a data which is likely to improve the classifier. Another issue in this problem is difficulty of the estimation of the accuracy of the obtained classifier on the remaining data. Because we choose data likely to be a target data in the first phase, our labeled data set obtained in the first phase is biased towards the target class. On the contrary, the remaining data is biased towards the non-target classes. In order to know on what point we can trust the classifier and switch from human annotation to the obtained classifier, we need to estimate the accuracy of the current classifier on the remaining biased data only by using the currently available labeled data, which is biased in the opposite way.

We call this problem target-extraction-learning problem. An example scenario of a potential application is as follows. Suppose we have a natural disaster, and we want to collect tweets on Twitter asking for rescue as fast as possible. Because we cannot examine all the tweets, we want to train a classifier, but the characteristics of tweets asking rescue is different from disaster to disaster, and we do not have any labeled data sets or pretrained classifiers. We need to label the data first, but it is waste of time to randomly sample tweets from Twitter and ask the annotators to label them. We should first examine tweets that are the most likely to be relevant, and if it is asking for rescue, we should send the rescue as soon as possible. We repeat it, and after obtaining enough labeled data, we should switch to the prediction by the classifier because human judgement takes longer time than the classifier for a huge number of tweets. Another example scenario is to find high risk patients among a big group of people without existing labeled data to train a predictive model, and under the constraint of a limited number of checkups we can perform [5].

In this paper, we introduce this new problem, and explain the difficulty of the problem by using our experimental results. We simulate the scenario with SVM implementation on document classification tasks and analyze the behavior of the classifier obtained at each step. We also discuss how to control the dilemma mentioned before, suggest a simple method of ballancing the dilemma explained above, and show its behavior in our simulation.

The remainder of the paper is organized as follows. In the following section, we first discuss several related literatures about active learning and also about learning-to-enumerate problem, which is a problem similar to our target-extraction-learning problem. We then define the details of our problem in Section 3, and propose an approach trying to provide possible solution to the problem in Section 4. In the experiment section, we simulate our framework on both balanced and unbalanced data, and report the results. Finally, we discuss

what we can conclude from the experiments and further issues to the mentioned problems.

2 Related Work

Two most important problems related to our problem is active learning and learning-to-enumerate. In this section, we briefly explain existing research on these two problems.

2.1 Active Learning [4]

As the boosting development of supervised machine learning, models are becoming more and more hungry for labeled training data. As we all know, human annotation may cost a huge amount of money and time. To handle this problem, many active learning algorithms have been developed. In the survey paper by Settles [4], the author summarized three typical categories of active learning algorithm: pool-based (the learner picks up a sample from the pool of unlabeled data by means of a querying function), stream-based (the learner tells whether to do annotation in a stream of unlabeled data) and membership queries (the learner generates data from the feature space of unlabeled data).

The problem in this research is close to the pool-based active learning, especially in the aspect of a fixed pool of unlabeled data which do not change during the annotation phase. On the other hand, typical active learning only focus on the performance of the model, while in our research, we try to ballance the performance of the model which is used in the second phase, and quick collection of target data in the first phase.

Another difference between the ordinary problem setting in active learning and our problem setting is the assumption on the properties of the remaining data. In the ordinary active learning, while we assume a closed data set when we select samples to label, we assume an open data set when we measure the performance of the obtained model. We assume that the selection of samples to label do not affect the properties of the remaining data. On the other hand, in our problem, we assume that the selection of samples in preference to the target data results in the remaining data biased towards the non-target data.

2.2 Learning to Enumerate [3]

Sometimes we want to find data satisfying some particular conditions and our goal is to find such data from the whole dataset as fast as possible. One scenario might be to find tweets asking for rescue after some natural disaster, explained before.

In such a case, if we have an extensive labeled dataset to train our model in advance, the best strategy would be to train a model with it and to solely rely on the predictions by the model to identify all data of the target class. However, when we do not have any labeled data in advance, one

approach is to construct a training set while examining data that are most likely to be the target data one by one, and use the current trained model to choose the next candidate. In other words, we can continue the first phase of our method until we find all the target data. In this setting, we have a dilemma between querying a label for a data which is likely to be a target data and querying for a data which is likely to improve the classifier, as explained before. Jörger et al. [3] call this problem the learning-to-enumerate problem.

This is a typical exploitation vs. exploration dilemma that has been studied extensively in reinforcement learning. Jörger et al. [3] proposed a simple ϵ -greedy like strategy, and tested it with different base learners and heuristic functions, on 19 small and medium-sized public datasets accessible through the UCI Machine Learning Repository. The results of their experiments shows that the best result was achieved by an exploitation only ($\epsilon = 0$) strategy. Here the exploitation means the model picks up a data that is the most likely to be of the target class, i.e., the data for which the model has the highest confidence.

However, there is a difference in our research from their learning-to-enumerate problem. Their purpose is to find all target data by human labeling, while in our research, we want to switch to the prediction by the model at certain point, when we can be sure that the model is well trained.

3 Problem Definition

In target-extraction-learning problem, as the literal meaning, we want to extract target data faster, and also try to learn a model. We presume that there is a fixed pool of unlabeled data, and the final purpose is to find target samples using as few non-target annotation as possible. To achieve this goal, we consider two-phase method. At the early phase, the system choose data that are likely to be target data, query for human annotation, and train a model with labeled data. At the later phase, we switch from human annotation to the prediction by the model, and provide a ranking of data based on the likelihood of being a target data.

At the early phase, it is reasonable just to make use of the learning-to-enumerate algorithm, but the point is that, the labeled data, which is used to train a machine learning model, is biased towards target samples, so we can not estimate how the model will perform on the unlabeled data. In order to know the details, we build a framework to simulate this annotation process.

A flow chart of the simulation process is shown in Figure 1. The framework works as follows:

- Iteratively pick up a sample from the unlabeled data pool at each step according to a certain policy (similar

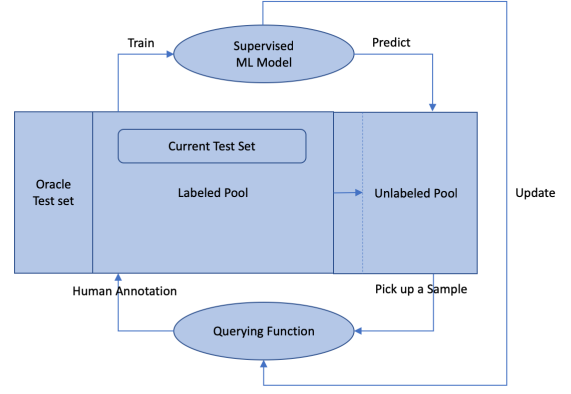


Figure 1: Flow Chart of Simulation Process

to querying function in active learning).

- We usually use all labeled data to train the model, but in this framework, we split a small part of labeled pool as “current test set”, to evaluate the model at each step. Metrics, such as precision and recall, measured on this current test set are available in real application.
- In comparison, we provide a test set with label to evaluate the model without bias, which we call it “oracle test set”. Various metrics measured on the oracle test set are only available in this simulation for research purpose, and cannot be known in real application.
- Use the labeled data pool to train a machine learning model, and use the model to update the querying function if needed.

3.1 Evaluation

In order to see every detail in the annotation simulation, we evaluate several metrics. Note that there are metrics that we can obtain in real application, and metrics that we cannot obtain in real application. In the list below, for the metrics that can be known in real application, the corresponding checkbox is checked. The others are used only for the evaluation in this simulation.

- ☐ Accuracy, recall, precision and f1-score on oracle test set.
- ☒ Accuracy, recall, precision and f1-score on current test set.
- ☐ Positive coverage: the ratio of retrieved target samples to all target samples.
- ☒ Recent labeled target proportion: the proportion of retrieved target samples in recent K labeled samples.
- ☐ R-precision on unlabeled data.

We need an explanation on the last one in the list above, i.e., R-precision. We want the model to be able to give a

target likelihood ranking of the remaining unlabeled data. To measure that ability of the model, we use R-precision. R-precision on the remaining data is defined as follows:

$$\text{R-precision} = \frac{A}{R}$$

where R is the number of target data in the remaining data, and A is the number of target data ranked within top- R by the model. However, if we use the ordinary R-precision to measure the performance of the model on the remaining data, the comparison can be unfair. For example, if we have been very successful in collecting target data and only a small number of difficult-to-find samples are left in the remaining data, it is difficult to achieve high R-precision on that remaining data. On the contrary, if we have been unsuccessful and many target data is still left in the remaining data, it is easier to achieve high R-precision on the remaining data. To avoid this unfairness, we add the number of already-found target samples to both numerator and denominator of R-precision. That is, we use R-precision defined by the formula below:

$$\text{R-precision} = \frac{A + F}{R + F}$$

where F is the number of already-found target data.

Metrics on oracle test set are used to give the “real” performance of obtained model in this simulation, and metrics on current test set are used to estimate the performance of the model on the remaining data in real applications (since oracle test set is not available there). At the early phase, we evaluate the performance of finding target data by positive coverage. At the later phase, we use R-precision defined above to evaluate the target likelihood ranking of unlabeled data predicted by the obtained model.

4 Proposed Approaches

In target-extraction-learning, we have a dilemma that it is hard to let these two purposes (improving the model and finding target data fast) be fulfilled at the same time. The reason is that if we find target data faster, the data used to train the model becomes biased towards target data, and as a result, the performance of the model will be deteriorated compared with the training with random sampling.

To handle this dilemma, we propose a compromised approach that uses a “warm” start with random sampling to balance the trade-off. We use random sampling only at the very beginning of the first phase (for example, 1/10 of the whole unlabeled pool), and switch to target-extraction after that. The point is that, we can make use of random sampling phase to estimate the distribution of target data in the unlabeled pool.

Another issue in our problem is the difficulty of estimating

the accuracy of the obtained model on the remaining data due to the bias in the training data set and the remaining data set. To estimate the performance of obtained model, we observe the curves of metrics on oracle test set and current test set, and have found a regular phenomena of “current precision pit”, which is able to indicate the relation between oracle and current test set. We will explain the details in the next section reporting our experimental results.

5 Experiments

In experiments, we run simulations both on balanced data and unbalanced data, and compared our methods with random sampling. The data sets we used in our experiments are as follows:

- Dataset 1: Text message spam classification, 1500 messages, 50% target data (spam).
- Dataset 2: Movie review sentiment classification, 4000 reviews, 20% target data (positive reviews).

For machine learning model, we chose Support Vector Machine (SVM). We have also compared Long Short Term Memory (LSTM) and Random Forest (RF). The results showed that the linear SVM and LSTM give the similar performance, while SVM is far more efficient to train. RF gives poor performance without hyper-parameter tuning, but we cannot tune hyper-parameters in real application, so we finally chose linear SVM with all default hyper-parameters in scikit-learn library [6]. In addition, we use class weight in training, as the data could be unbalanced.

On oracle test set, we use 5-fold cross-validation in order to obtain stable results. On current test set, we do not use cross-validation due to efficiency reason.

Figure 2 to 11 show the metrics observed in our simulations. The horizontal axis represents the size of the labeled pool, also can be considered as “steps”. We use the learning-to-enumerate algorithm [3] as pure target-extraction baseline, and also use the method that combine it with 10% random sampling warm start. At each step, random sampling picks up a sample randomly, while target-extraction picks up a sample which has the largest distance to the decision boundary on the target side, in contrast to uncertainty sampling, which picks up a sample closest to the boundary.

5.1 Results on Dataset 1

In Figure 4 showing positive coverage, we can see that target-extraction works expectedly. On the other hand, in Figure 2, since both accuracy and f1-score on oracle test set show good results at the final step, we can assert that the model is well trained. However, random sampling gives a much better learning speed than target-extraction, since the




Figure 2: Metrics on Oracle Test Set

labeled data does not have any bias. On the other hand, in target-extraction, the model does not learn well at the early phase, since the labeled data used to train is biased towards target samples. As the model has found most of target samples (around the step 600), non-target samples begin to increase in the labeled data, and the model starts to learn the classification task properly. In the precision on the current test set (Figure 3), there is an obvious pit around the step 700. We call it “current precision pit”. The reason of this pit is as follows. At the early phase, the training data set is biased toward the target data, and the model does not well trained. However, the current test set is also biased toward the target data, so the model shows good performance on the current test set, and we cannot know that the model is not well-trained only from the current test set metrics. At the later phase, as the current test set gradually becomes balanced, these metrics begin to go down, but when the model begins to learn the classification task properly, the metrics go up again.

In Figure 5 showing R-precision, we can see that the model shows poor R-precision compared with random sampling in early steps, when the model is not well trained. On the contrary, when the model starts to perform well around step 700, R-precision of target-extraction surpass that of random

sampling, which indicates that the model trained by target-extraction algorithm also provides a better ranking. This satisfies our second purpose of providing a good target likelihood ranking, so that we can switch from human annotation to the prediction by the model.

The warm start method has shown a compromised result. Although it do not extract target data as fast as pure target-extraction method, it was able to provide a much higher R-precision.

Figure 6, which shows the recent labeled target proportion, indicates the proportion of target samples in the recent ten labeled data. Random sampling shows a relatively stable curve as expected, while target-extraction picks much more target samples at early steps, but the proportion goes down at later steps when less target samples remain in unlabeled data. Since this is a metric which can be known in real application, it provide a possibility to estimate the target proportion of unlabeled data. Also, in warm start method, it provided a priori of the whole unlabeled pool as expected.

5.2 Results on Dataset 2

Figure 7 to 11 show the result for the second data set. In Figure 8, although the current precision pit came earlier and lasted for longer than in the previous experiment with the balanced dataset, f1-scores on oracle test set and current test

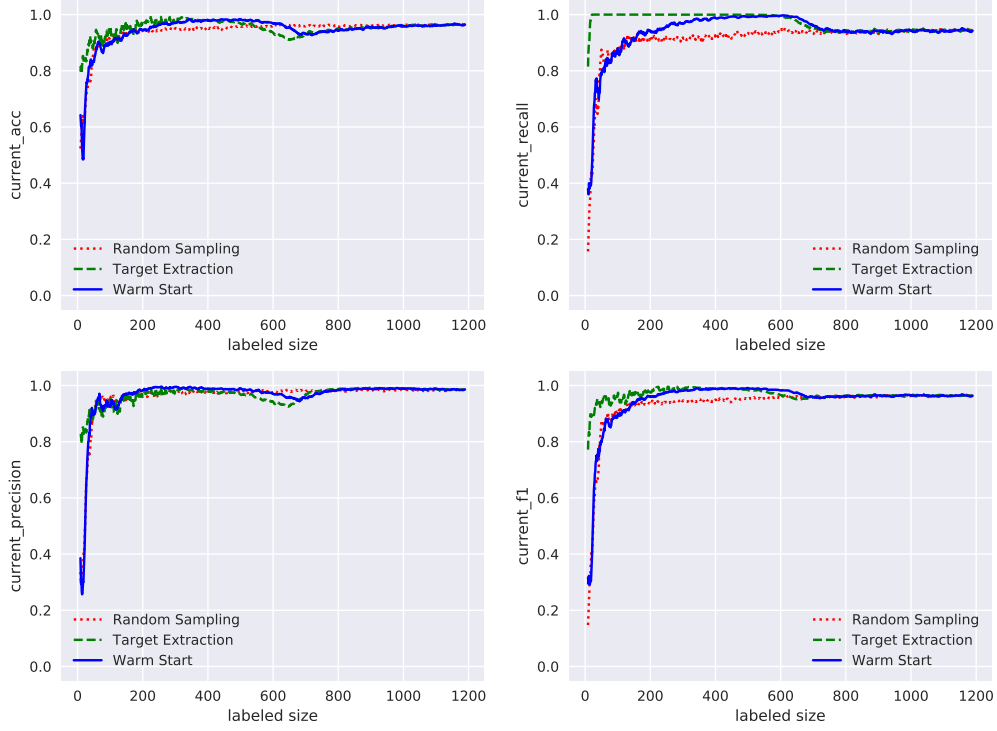


Figure 3: Metrics on Current Test Set




Figure 4: Positive Coverage




Figure 6: Recent Target Proportion




Figure 5: R-Precision

set still converged, which means we can use current performance to estimate oracle performance after the pit appears.

The f1-score on oracle test set of Dataset 1 (Figure 7) shows that the model with target-extraction method learned much slower than with random sampling, due to the bias towards target data. However, in unbalanced dataset, model with target-extraction method learned much faster than random sampling. This is because the model tried to fetch target data, which is minority in unlabeled data pool, so that the model was able to learn on a more balanced dataset than random sampling. In this way, target-extraction method not only fetched target data faster, but contributed to model

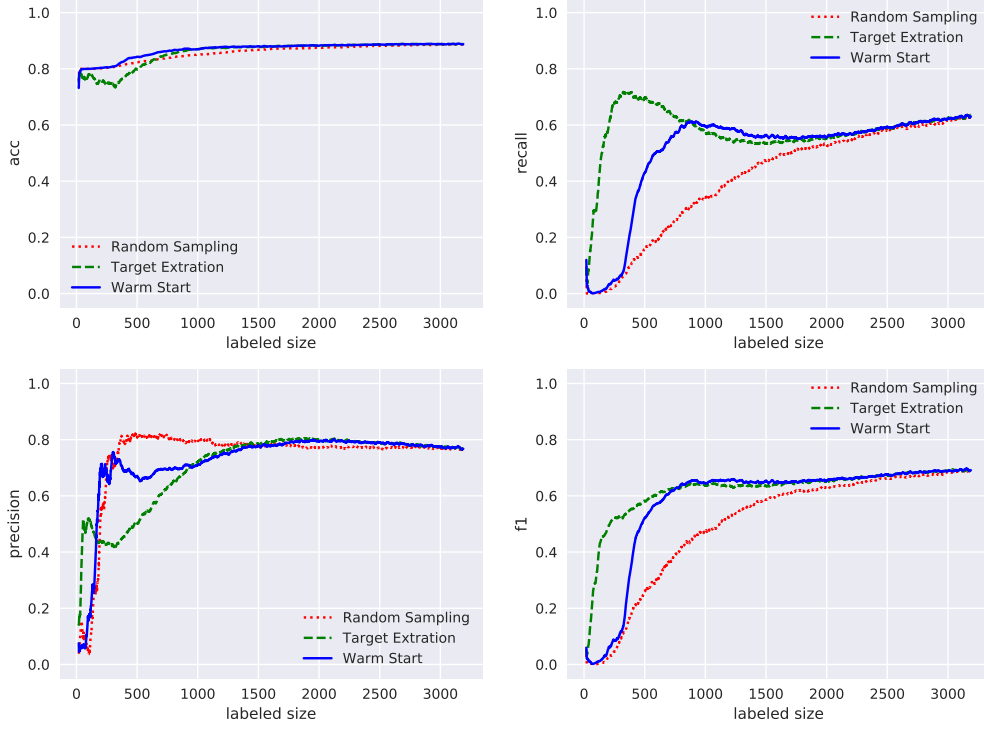


Figure 7: Metrics on Oracle Test Set

performance as well for this data set where the target data is minority. Hence, we can expect that if the target data is majority in the unlabeled pool, the model will learn on a even more unbalanced dataset, and the performance will go worse than random sampling.

Figure 10 shows that random sampling and target-extraction on this unbalanced Dataset 2 gave the results opposite to those on balanced Dataset 1. It is easy to understand because a well-learned model should provide a better ranking of target likelihood. We can also conclude that the warm start method gives a compromised performance of R-precision on both balanced and unbalanced datasets.

6 Conclusion

We defined a new problem called target-extraction-learning. The purpose is to extract target data faster, and we propose a two-phase approach. At the early phase, the system queries samples most likely to be a target data for annotation, and at the later phase, we switch from human annotation to the prediction by the model. We also analyze the difficulty in the model performance estimation due to the bias towards target data in the training data set. To handle this difficulty, we use a phenomena we found in our experiments. We can observe a “current precision pit” in

precision curve on current test set, and we are able to estimate the model performance since the results on current test set and oracle test set begin to converge after striding over the pit. Finally, to handle the dilemma between model performance and target extraction, we proposed a compromised approach with a warm start of random sampling, and it gives compromised results, both on balanced and unbalanced data.

However, the “current precision pit” can only be observed after the model is well-trained. If the model cannot be well trained on a certain dataset, the pit will not appear. We still need a way to estimate the performance even in such cases. In addition, estimation using the experimental phenomena is still ambiguous, and it is better to give more statistical evidence, such as the expectation and variance to the performance of current model.

7 Acknowledgment

This work was supported by JST CREST (JPMJCR16E3), and JSPS KAKENHI Grant Numbers 18H03245, 16K12430, and 18K11425.

References

- [1] Patrick Jörger, Yukino Baba, and Hisashi Kashima. Learning to enumerate. *Artificial Neural Networks and Machine*

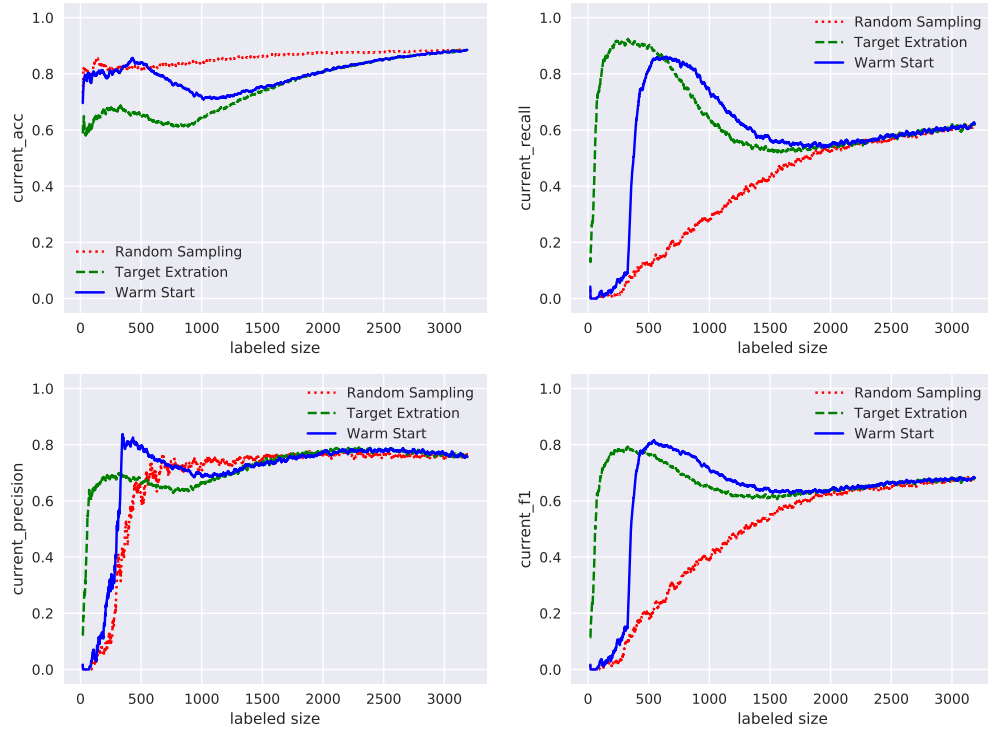


Figure 8: Metrics on Current Test Set




Figure 9: Positive Coverage




Figure 11: Recent Target Proporttion




Figure 10: R-Precision

- Learning*, 2016.
- [2] Y. Baba et al. Predictive approaches for low-cost preventive medicine program in developing countries. *Proc. of KDD*, pages 1681–1690, 2015.
 - [3] Anita Krishnakumar. Active learning literature survey. 2007.
 - [4] David D. Lewis and William A. Gale. A sequential algorithm for training text classifiers. *Proc. of SIGIR*, 1994.
 - [5] F. Pedregosa et al. Scikit-learn: Machine learning in python. *Journal of Machine Learning Research*, 12:2825–2830, 2011.